

Reasoning with Happens-Before Relations about Concurrent Programs in the Theta Framework

Short toolpaper

Csanad Telbisz, Levente Bajczi, Danel Szekeres,
Andras Voros, Istvan Majzik

December 7, 2025



Agenda

Introduction

Context: what is the problem?
What do we do? **Why** do we care?

Agenda

Introduction

Context: what is the problem?
What do we do? **Why** do we care?

Verification approaches (existing & novel)
Optimization techniques (novel)

Techniques

Agenda

Introduction

Context: what is the problem?
What do we do? **Why** do we care?

Verification approaches (existing & novel)
Optimization techniques (novel)

Techniques

Evaluation

Experiment and data analysis
Evaluation of impact

Agenda

Introduction

Context: what is the problem?
What do we do? **Why do we care?**

Verification approaches (existing & novel)
Optimization techniques (novel)

Techniques

Evaluation

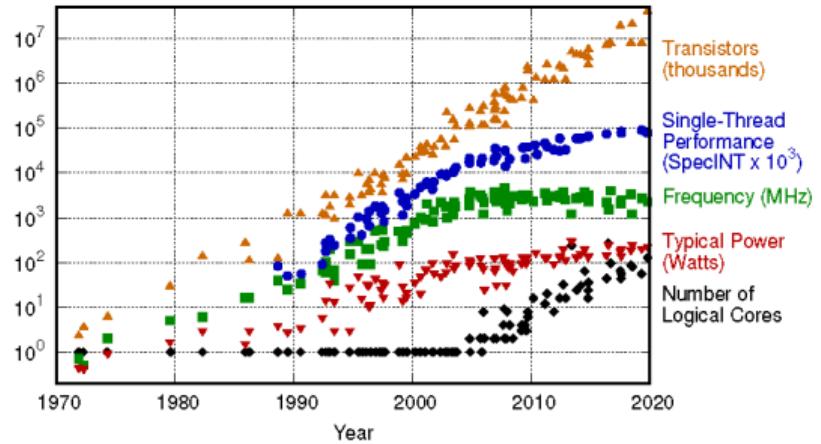
Experiment and data analysis
Evaluation of impact

Theta as a Verifier
Theta as a Framework

Availability

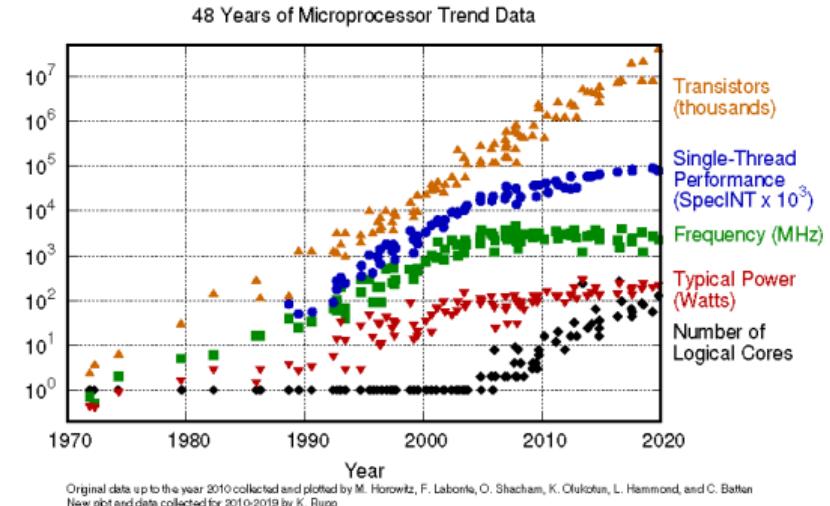
Context

48 Years of Microprocessor Trend Data



Context

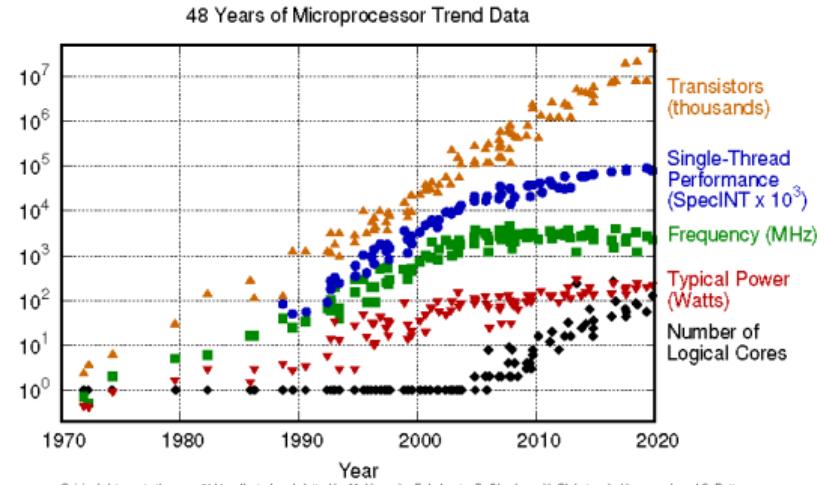
- ▶ Safety critical systems need **multicore** CPUs
- ▶ Further performance: **relaxed** memory models



Context

- ▶ Safety critical systems need **multicore** CPUs
- ▶ Further performance: **relaxed** memory models

We need new
formal techniques!

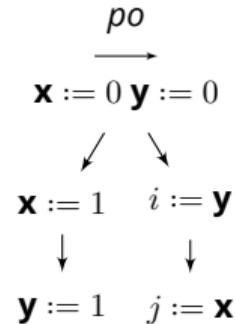


Introduction

- ▶ Emerging method for verifying concurrency: **Happens-before relation**
 - ▶ Decouples data- and control-flow
 - ▶ Enables semi-modular verification

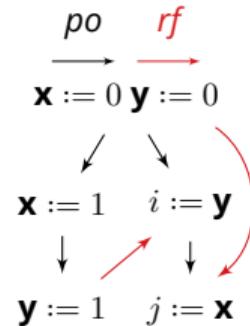
Introduction

- ▶ Emerging method for verifying concurrency: **Happens-before relation**
 - ▶ Decouples data- and control-flow
 - ▶ Enables semi-modular verification



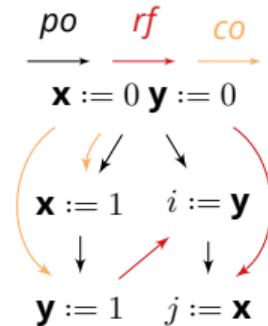
Introduction

- ▶ Emerging method for verifying concurrency: **Happens-before relation**
 - ▶ Decouples data- and control-flow
 - ▶ Enables semi-modular verification



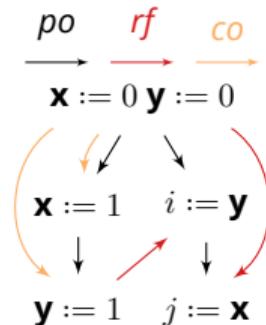
Introduction

- ▶ Emerging method for verifying concurrency: **Happens-before relation**
 - ▶ Decouples data- and control-flow
 - ▶ Enables semi-modular verification



Introduction

- ▶ Emerging method for verifying concurrency: **Happens-before relation**
 - ▶ Decouples data- and control-flow
 - ▶ Enables semi-modular verification



- ▶ Causality and memory semantics: happens-before (*hb*-) order
- ▶ **Verification idea:** encode threads *and* happens-before constraints
- ▶ **Error property:** reachable thread state, and \exists consistent *hb*-order

Verification Approach (Overview)

Given a program with...

POSIX/C11 threads
Nondeterministic inputs
Thread-local assertions

Verification Approach (Overview)

Given a program with...

POSIX/C11 threads
Nondeterministic inputs
Thread-local assertions

Is any assertion violated?

Verification Approach (Overview)

Given a program with...

POSIX/C11 threads
Nondeterministic inputs
Thread-local assertions

Is any assertion violated?

initially: $x_0 = y_0 = 0$

Thread t_1 Thread t_2

```
if ( $y_1 == 1$ )  $x_1 = 1$ ;  $y_3 = x_3$ ;  
else  $x_2 = 1 - y_2$ ;  $x_4 = 1 - y_4$ ;
```

finally: assert($x_5 == 1 \parallel y_5 == 1$)

Verification Approach (Overview)

Given a program with...

POSIX/C11 threads
Nondeterministic inputs
Thread-local assertions

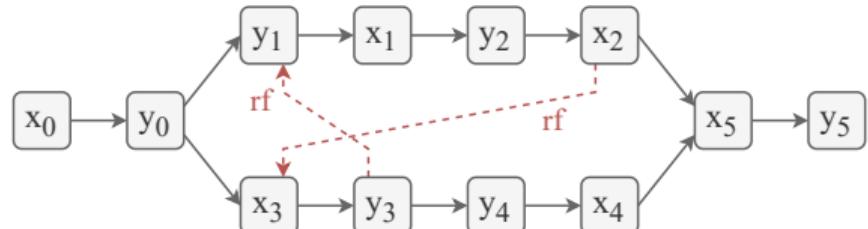
Is any assertion violated?

initially: $x_0 = y_0 = 0$

Thread t_1 Thread t_2

```
if(y1==1) x1=1; y3=x3;
else x2=1-y2; x4=1-y4
```

```
finally: assert(x5==1 || y5==1)
```



Verification Approach (Overview)

Given a program with...

POSIX/C11 threads
Nondeterministic inputs
Thread-local assertions

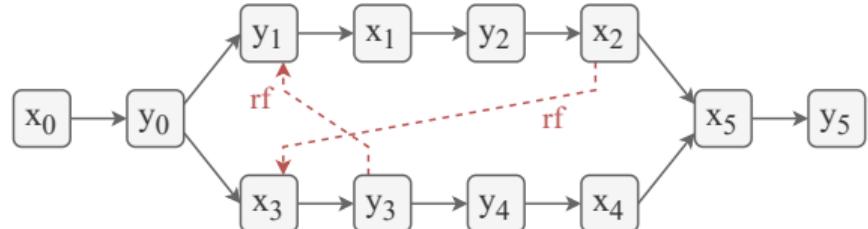
Is any assertion violated?

initially: $x_0 = y_0 = 0$

Thread t_1 Thread t_2

```
if( $y_1 == 1$ )  $x_1 = 1$ ;  $y_3 = x_3$ ;  
else  $x_2 = 1 - y_2$ ;  $x_4 = 1 - y_4$ ;
```

finally: assert($x_5 == 1 \parallel y_5 == 1$)



Unroll

Verification Approach (Overview)

Given a program with...

POSIX/C11 threads
Nondeterministic inputs
Thread-local assertions

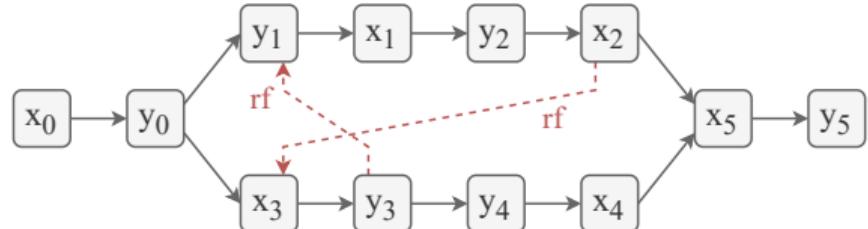
Is any assertion violated?

initially: $x_0 = y_0 = 0$

Thread t_1 Thread t_2

```
if( $y_1 == 1$ )  $x_1 = 1$ ;  $y_3 = x_3$ ;  
else  $x_2 = 1 - y_2$ ;  $x_4 = 1 - y_4$ ;
```

finally: assert($x_5 == 1 \parallel y_5 == 1$)



Unroll

→ CSSA

Verification Approach (Overview)

Given a program with...

POSIX/C11 threads
Nondeterministic inputs
Thread-local assertions

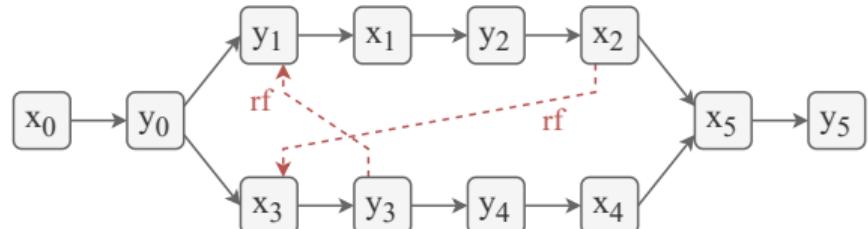
Is any assertion violated?

initially: $x_0 = y_0 = 0$

Thread t_1 Thread t_2

```
if( $y_1 == 1$ )  $x_1 = 1$ ;  $y_3 = x_3$ ;  
else  $x_2 = 1 - y_2$ ;  $x_4 = 1 - y_4$ ;
```

finally: assert($x_5 == 1 \parallel y_5 == 1$)



Unroll

CSSA

SMT

Verification Approach (Overview)

Given a program with...

POSIX/C11 threads
Nondeterministic inputs
Thread-local assertions

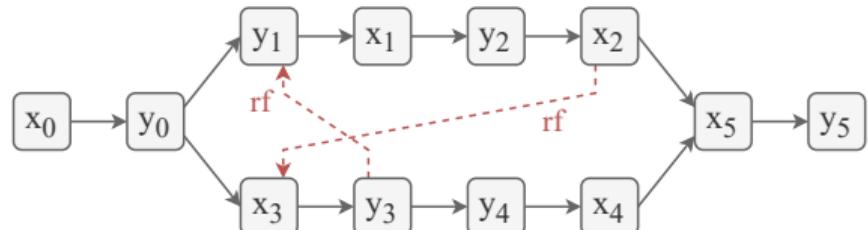
Is any assertion violated?

initially: $x_0 = y_0 = 0$

Thread t_1 Thread t_2

```
if(y1 == 1) x1=1; y3=x3;  
else x2=1-y2; x4=1-y4;
```

finally: assert($x_5 == 1 \ || \ y_5 == 1$)



Unroll

CSSA

SMT

Solve

Verification Approach (Overview)

Given a program with...

POSIX/C11 threads
Nondeterministic inputs
Thread-local assertions

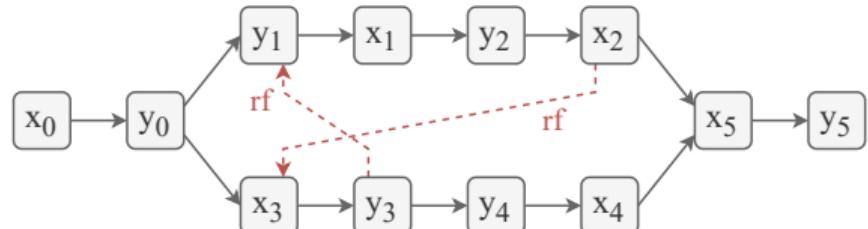
Is any assertion violated?

initially: $x_0 = y_0 = 0$

Thread t_1 Thread t_2

```
if(y1 == 1) x1=1; y3=x3;  
else x2=1-y2; x4=1-y4;
```

finally: assert($x_5 == 1 \parallel y_5 == 1$)



Unroll

CSSA

SMT

Solve

Verdict

Verification Approach (Overview)

Given a program with...

POSIX/C11 threads
Nondeterministic inputs
Thread-local assertions

Is any assertion violated?

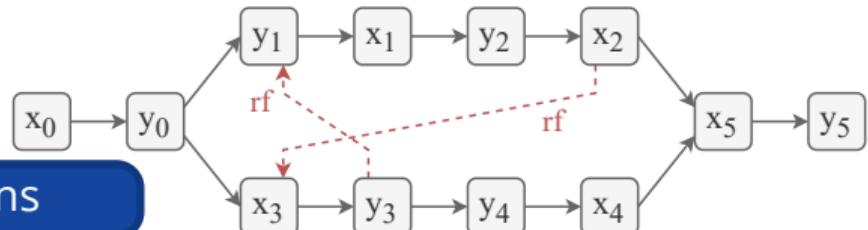
initially: $x_0 = y_0 = 0$

Thread t_1 Thread t_2

```
if(y1 == 1) x1=1; y3=x3;  
else x2=1-y2;
```

finally: assert($x_5 = y_5$)

Optimizations



Unroll

CSSA

SMT

Solve

Verdict

Verification Approach (Overview)

Given a program with...

POSIX/C11 threads
Nondeterministic inputs
Thread-local assertions

Is any assertion violated?

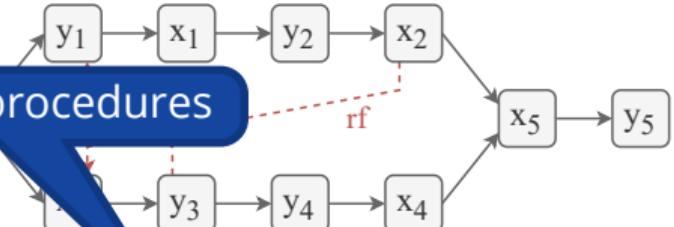
initially: $x_0 = y_0 = 0$

Thread t_1 Thread t_2

```
if(y1 == 1) x1=1; y3=x3;  
else x2=1-y2; y1=1-y1.  
finally: assert(x5=y5)
```

Decision procedures

Optimizations



Unroll

CSSA

SMT

Solve

Verdict

Verification Approach (Overview)

Given a program with...

POSIX/C11 threads
Nondeterministic inputs
Thread-local assertions

Is any assertion violated?

initially: $x_0 = y_0 = 0$

Thread t_1

Thread t_2

```
if(y1 == 1) x1=1; y3=x3;  
else x2=1-y2; y5=1-y5;  
finally: assert(x5=y5)
```

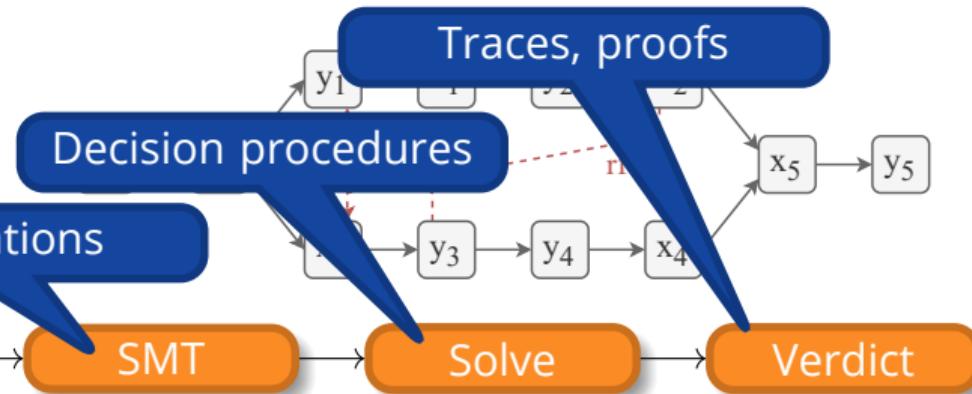
Unroll

CSSA

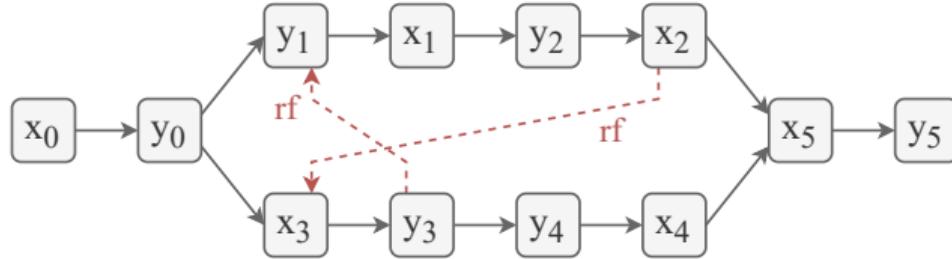
SMT

Solve

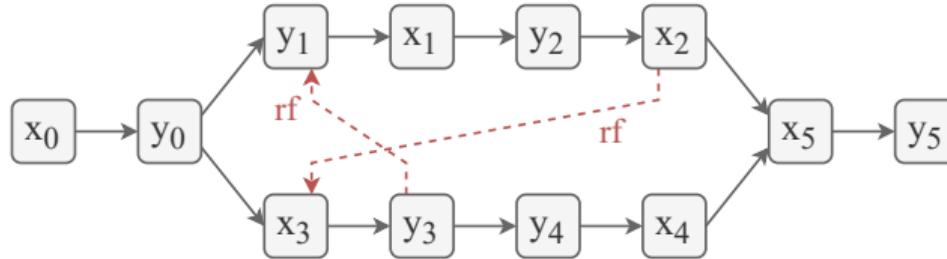
Verdict



Techniques - Integer Difference Logic

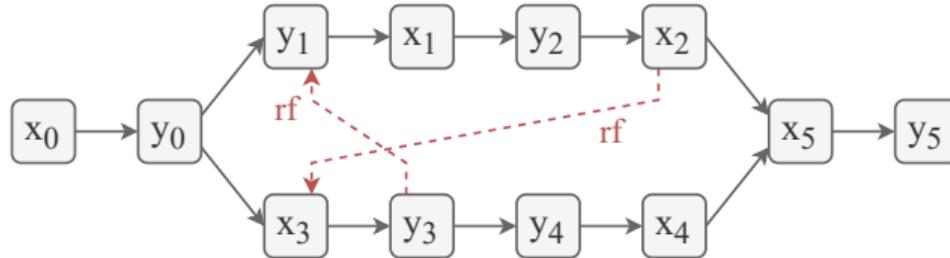


Techniques - Integer Difference Logic



- ▶ Each event has a *clock* integer variable
- ▶ A happens-before order implies an integer order

Techniques - Integer Difference Logic

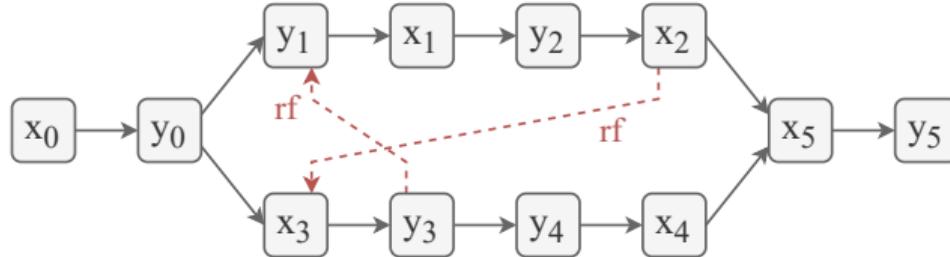


- ▶ Each event has a *clock* integer variable
- ▶ A happens-before order implies an integer order

Advantages

Easy encoding, any solver can be used

Techniques - Integer Difference Logic



- ▶ Each event has a *clock integer variable*
- ▶ A happens-before order implies an integer order

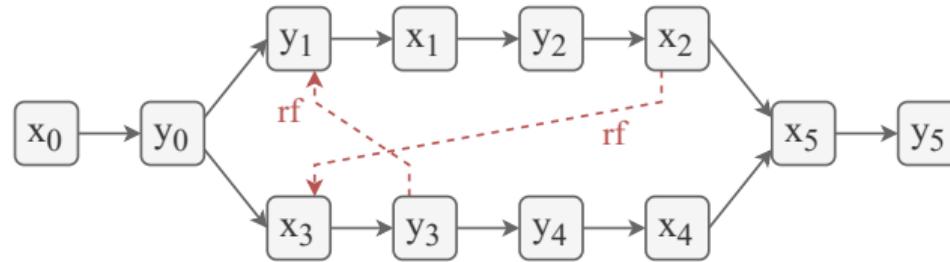
Advantages

Easy encoding, any solver can be used

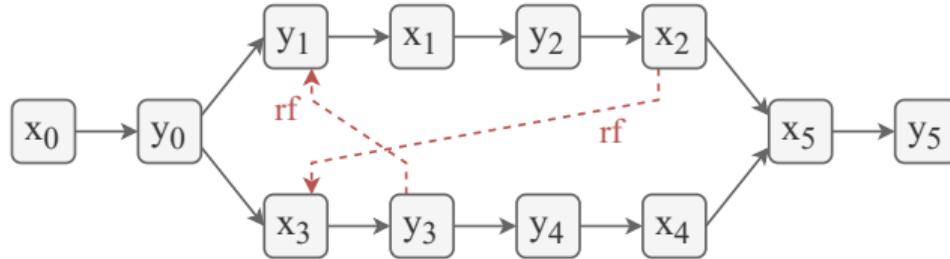
Disadvantages

Overly specific, thus suboptimal

Techniques - Refinement Step-by-Step

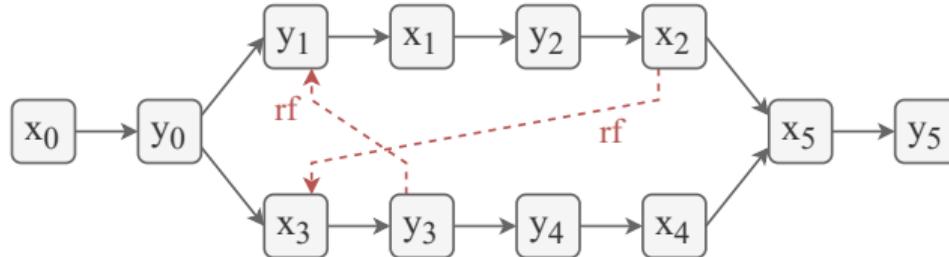


Techniques - Refinement Step-by-Step



- ▶ We track *happens-before* orders externally to the solver
- ▶ Any model is validated, and potentially *refined*
- ▶ The SMT solver does not propagate, we do

Techniques - Refinement Step-by-Step

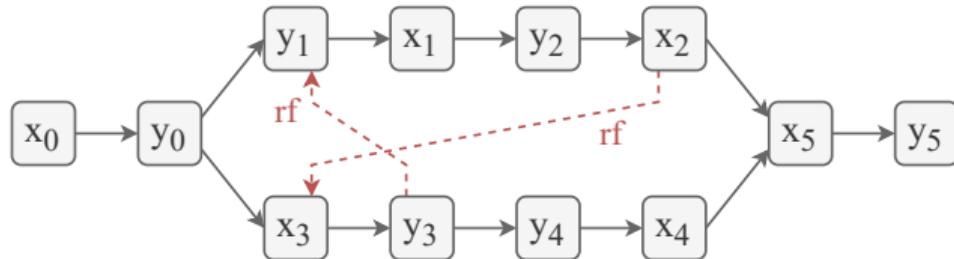


- ▶ We track *happens-before* orders externally to the solver
- ▶ Any model is validated, and potentially *refined*
- ▶ The SMT solver does not propagate, we do

Advantages

Any solver can be used

Techniques - Refinement Step-by-Step



- ▶ We track *happens-before* orders externally to the solver
- ▶ Any model is validated, and potentially *refined*
- ▶ The SMT solver does not propagate, we do

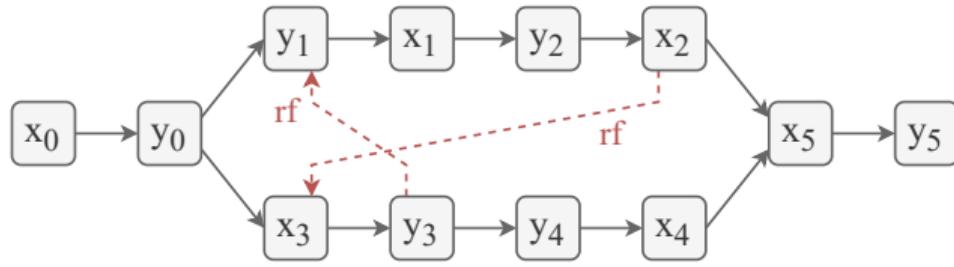
Advantages

Any solver can be used

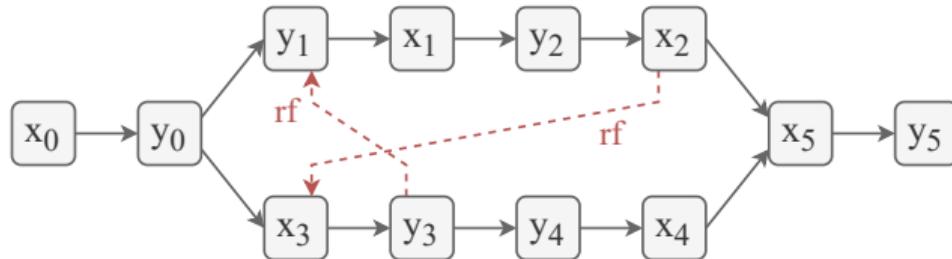
Disadvantages

Number of SMT queries increase

Techniques - User Propagator

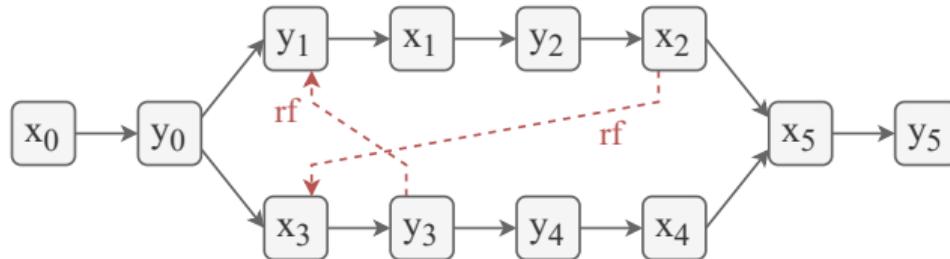


Techniques - User Propagator



- ▶ We track *happens-before* orders internally in the solver
- ▶ Derived relations and conflicts are propagated
- ▶ A model is always consistent with *hb* acyclicity rules

Techniques - User Propagator

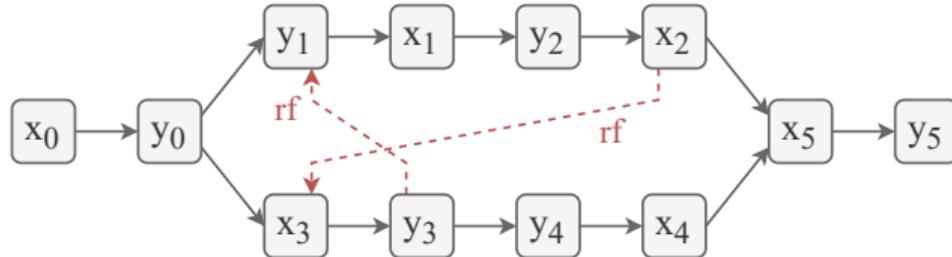


- ▶ We track *happens-before* orders internally in the solver
 - ▶ Derived relations and conflicts are propagated
 - ▶ A model is always consistent with *hb* acyclicity rules

Advantages

Native solution

Techniques - User Propagator



- ▶ We track *happens-before* orders internally in the solver
- ▶ Derived relations and conflicts are propagated
- ▶ A model is always consistent with *hb* acyclicity rules

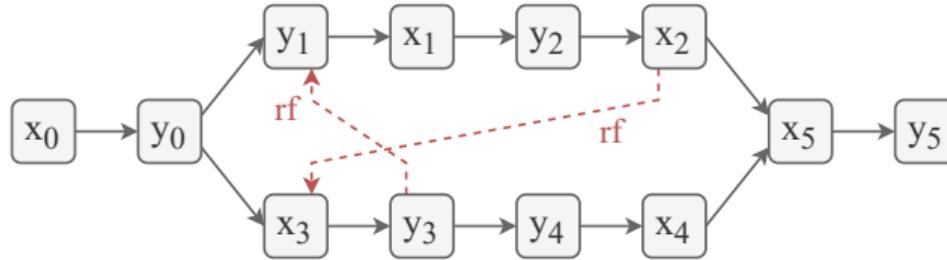
Advantages

Native solution

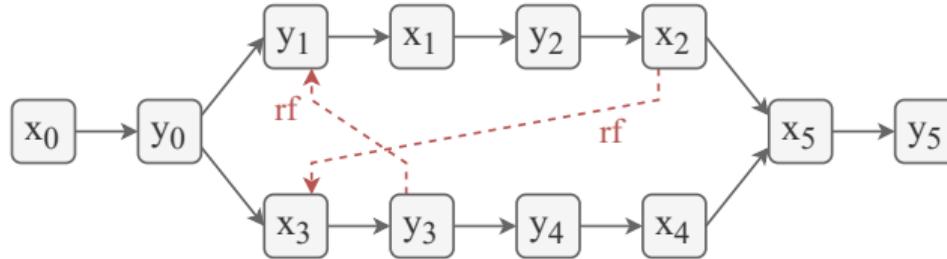
Disadvantages

Only certain SMT solvers are usable

Optimization: Automatic Conflict Avoidance



Optimization: Automatic Conflict Avoidance



- ▶ **Idea:** we enumerate some potential cycles explicitly
- ▶ This is sound (we don't overconstrain the state space)
- ▶ This helps the solvers (based on empirical results)

Experimental Evaluation

Research Question 1

How do the decision procedures compare?

Research Question 2

How much does the optimization help?

Research Question 3

How does Theta compare to other tools?

Research Question 4

Does solver choice influence performance?

Results - RQ1

Research Question 1

How do the decision procedures compare?

Results - RQ1

Research Question 1

How do the decision procedures compare?

- ▶ SV-Benchmarks tasks: 725 concurrent C programs
- ▶ 15 CPU-minute timeout, 15GB RAM

Results - RQ1

Research Question 1

How do the decision procedures compare?

- ▶ SV-Benchmarks tasks: 725 concurrent C programs
- ▶ 15 CPU-minute timeout, 15GB RAM

Solved tasks

	IDL	RFN	PROP
Solved	398	409	410
Time (s)	26000	4150	5770

Results - RQ1

Research Question 1

How do the decision procedures compare?

- ▶ SV-Benchmarks tasks: 725 concurrent C programs
- ▶ 15 CPU-minute timeout, 15GB RAM

Solved tasks

	IDL	RFN	PROP
Solved	398	409	410
Time (s)	26000	4150	5770

- ▶ Integer difference logic is worse and slowest
- ▶ Step-by-step refinement is quickest and almost best
- ▶ User propagation is best but not quickest

Results - RQ2

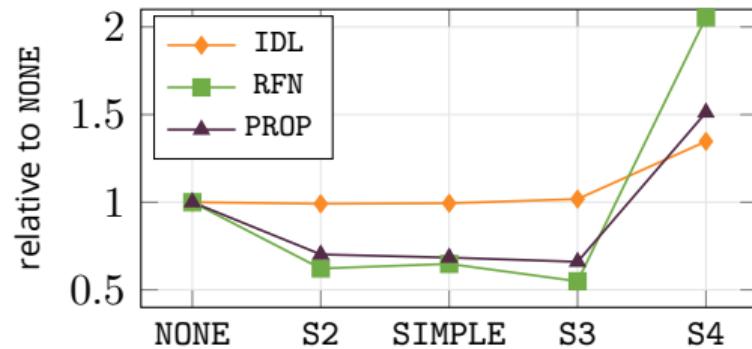
Research Question 2

How much does the optimization help?

Results - RQ2

Research Question 2

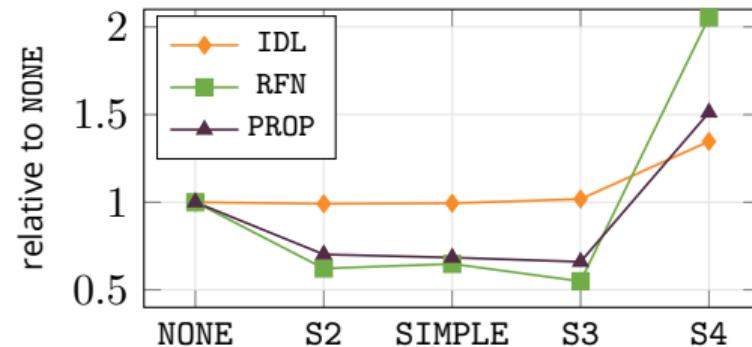
How much does the optimization help?



Results - RQ2

Research Question 2

How much does the optimization help?



- ▶ SN : Length of cycle hardcoded in SMT
- ▶ "Simple": 2, plus one derivation rule

Results - RQ3: Dartagnan

Research Question 3

How does Theta compare to other tools?

Results - RQ3: Dartagnan

Research Question 3

How does Theta compare to other tools?

	Theta (complete)			Dartagnan	
	IDL	RFN	PROP	Eager	Lazy
Solved/	398	409	410	456	457
filtered	398	409	410	434	433
Time (s)	26000	4150	5770	11200	6370

Results - RQ3: Deagle

Research Question 3

How does Theta compare to other tools?

	Theta (bounded)			Deagle
	IDL	RFN	PROP	
Solved/	538	554	553	623
filtered	538	554	553	577
Time (s)	35000	6020	8690	2020

- ▶ Bounded safety: unsound results

Results - RQ4

Research Question 4

Does solver choice influence performance?

Results - RQ4

Research Question 4

Does solver choice influence performance?

	IDL	RFN
Z3	538	544
cvc5	537	529
MathSAT	307	528
Princess	280	184

- ▶ Few new tasks solved

- ▶ Model checking framework originally designed for verification using abstraction and refinement (CEGAR)
- ▶ **Modular, extensible architecture**, supporting different verification engines (CEGAR, BMC, etc.).
- ▶ Emphasis on **reusability** across domains (control-flow automata, transition systems, statecharts, CHCs, ...).
- ▶ Supports **reusability** through regular updates on maven
- ▶ Full verifier packages via Docker, and Zenodo releases

Summary

Introduction

Context: what is the problem?
What do we do? Why do we care?

Verification approaches (existing & novel)
Optimization techniques (novel)

Techniques

Evaluation

Experiment and data analysis
Evaluation of impact

Theta as a Verifier
Theta as a Framework

Availability